

Client-Server Programming in Java-RMI

<http://wiki.duke.edu/display/scsc>
scsc@duke.edu

John Pormann, Ph.D.
jbp1@duke.edu

Outline

- ✧ Client-Server Parallelism
- ✧ Java RMI
- ✧ Using RMI for Client-Server Interactions
- ✧ Building the code
- ✧ Launching the application
- ✧ Performance Issues/Hints

Client-Server Parallelism


- ✧ Multiple independent clients that ONLY interact with the server
 - ◆ Client-to-Client data exchange may not be needed at all, or else it is handled through the server as 2 separate exchanges
 - Client #1 talks to Server
 - Server talks to Client #2
 - ◆ Server coordinates any data distribution that is needed
- ✧ Sometimes referred to as “Remote Procedure Calls”
 - ◆ Clients interact with server as though it was a local subroutine:

```
my_data = get_job_data( server );  
value = compute( my_data );  
err = post_job_result( server, my_data, value );
```

 - Where the `get_job_data` and `post_job_results` functions just happen to execute on another machine
 - ◆ What if server has too many clients and slows down?

Java calls them “Remote Method Invocations”

Client-Server Parallelism, cont'd



Client-Server Parallelism, Examples

- ✧ File serving, Web serving
- ✧ Database query
 - ◆ Server hands out the data requested by each client
 - ◆ Presumably, each client knows to request different pieces
- ✧ Data-mining
 - ◆ Multiple clients scan database for different parameters
- ✧ Monte Carlo methods
 - ◆ Server ensures unique random number seeds for clients
- ✧ Function mapping, minimization, successive refinement
 - ◆ If $x=0$ to $x=0.3$ have “high” values, then the server can focus future computational efforts elsewhere
- ✧ Web services
 - ◆ A new/emerging model for distributed client-server applications

Client-Server Parallelism, More Ideas

- ✧ We often think of all clients as being “equal”
 - ◆ Each client runs the exact same program/same algorithm
 - ◆ This isn’t strictly necessary
 - Each client could run a different program/different algorithm
 - E.g. you might process the same image with 10 different DSP algorithms to keep the most accurate result
 - ▲ For convenience/sanity, it is often still good to have those 10 algorithms “wrapped” into 1 client-executable
- ✧ The batch scheduling system on the DSCR is a client-server app
 - ◆ So “Pool of Work” or “Job-level” parallelism can be done within a client-server framework

Pros/Cons of Client-Server Parallelism


Cons:

- ◆ You must be able to fit each client-job onto a single machines
 - Memory could be a limiting factor
- ◆ It will still take X hours for any one client to finish, it's just that you will get back 10 or 100 results every X hours
 - Sometimes (fast) partial results can be helpful
- ◆ Scalability is a function of the (coarse) problem "size"

Pros:

- ◆ It doesn't require "real" parallel programming!
- ◆ This can be one of the most efficient forms of parallelism
- ◆ It provides automatic load-balancing

Automatic Load-Balancing



The Java Programming Language

- ✧ Java is a full-featured, object-oriented programming language
 - ◆ “Write once, Run anywhere” -- uses a Java Virtual Machine to hide the details of the hardware
 - The Java Byte-Code (assembly language for the JVM) can be shipped to remote machines (clients) and you know it will execute there
 - ◆ “try..catch” blocks are particularly well-suited for distributed, networked programming (where things frequently fail)
- ✧ Sun pioneered the idea of “Remote Procedure Calls” (RPC), so it was natural for them to roll out “Remote Method Invocations” (RMI)
 - ◆ Same basic approach is used: automatically generates stubs for the remote methods, so the user only deals with the “real code” (not the network sockets, parallel programming code)

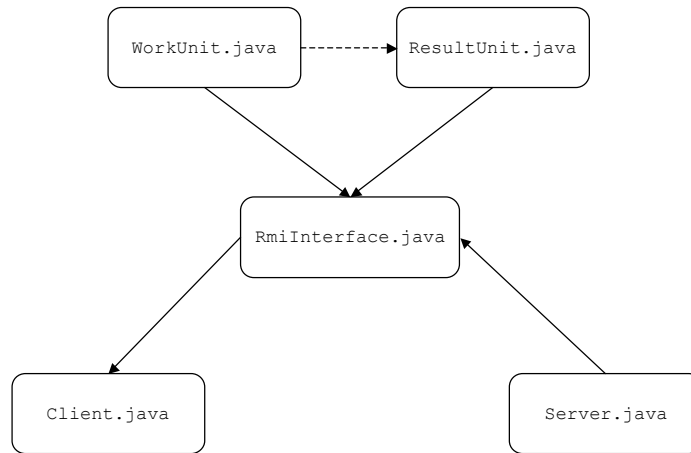
Java-RMI

- ✧ In Java, “Remote Methods” look like regular method (function) calls

```
RmiInterface server = (RmiInterface)Naming.lookup( some_url );  
WorkUnit wu = server.GetNextWorkUnit();
```

- ◆ The only hint that it is a remote method is the “Naming.lookup” call with a URL
 - (Hidden) Java RMI classes do all the “heavy lifting” -- opening network sockets, transfer the function arguments, retrieve the results, close the network sockets
- ◆ On the server side, there is a “Naming.rebind” call in the main method (to bind the server to a network location/URL)
 - Again, the Java RMI classes do the “heavy lifting” -- converting the remote method call into a “regular” method call on “normal” Java code running within the server

Java-RMI Classes



Java-RMI, Client Example

```

package ClientServer;
import java.rmi.*;
class Client {
    public static double DoCalculations( double alpha, double beta ) {
        double rtn;
        // do some highly complex calculations here
        rtn = (alpha+beta)*(alpha-beta);
        return( rtn );
    }

    public static void main( String[] argv ) {
        int SelfCID = Integer.parseInt(argv[0]);
        int NumClients = Integer.parseInt(argv[1]);
        String server_url = "/" + argv[2] + "/AlphaBetaServer";
        System.out.println( "Client "+SelfCID+" of "+NumClients+" starting" );
        try {
            RmiInterface server = (RmiInterface)Naming.lookup( server_url );
            for(int i=0;i<10;i++) {
                WorkUnit wu = server.getNextWorkUnit();
                double alpha = wu.alpha;
                double beta = wu.beta;
                System.out.println( "Client"+SelfCID+" : alpha="+alpha+" beta="+beta );
                double rtn = DoCalculations( alpha, beta );
                System.out.println( "Client"+SelfCID+" : rtb="+rtn );
                ResultUnit ru = new ResultUnit();
                ru.wu = wu;
                ru.rtnval = rtn;
                server.PutResultUnit( ru );
            }
        } catch (Exception e) {
            System.out.println ( "ComputeClient exception: " + e );
        }
    }
}
  
```

App-specific
start-up code

Java Classes
we provide

Remote Methods

Java-RMI, Server Example

```
package ClientServer;
import java.rmi.*;
import java.rmi.server.*;
class Server extends UnicastRemoteObject implements RmiInterface {
    private double next_alpha = 0.0;
    private double next_beta = 100.0;
    private double delta = 0.1;

    public Server( double a0, double b0, double dd ) throws RemoteException {
        next_alpha = a0;
        next_beta = b0;
        delta = dd;
    }
    public synchronized WorkUnit getNextWorkUnit() throws RemoteException {
        WorkUnit wu = new WorkUnit();
        wu.alpha = next_alpha;
        wu.beta = next_beta;
        next_alpha += delta;
        next_beta -= delta;
        return( wu );
    }
    public int PutResultUnit( ResultUnit ru ) throws RemoteException {
        // check that this is from a workunit we already processed
        if( ru.wu.alpha >= next_alpha ) {
            return( -1 );
        }
        // otherwise, we'll store the result somewhere
        System.out.println( "Result: a="+ru.wu.alpha+" b="+ru.wu.beta+" rtn="+ru.rtnval );
        return( 0 );
    }
    public static void main( String[] argv ) {
        try {
            Naming.rebind( "AlphaBetaServer", new Server(0.0,100.0,0.1) );
            System.out.println( "AlphaBetaServer is ready." );
        } catch (Exception e) {
            System.out.println( "AlphaBetaServer failed: " + e );
        }
    }
}
```

Java-RMI, Interface Example

```
public interface RmiInterface extends Remote {
    public WorkUnit getNextWorkUnit() throws RemoteException;
    public int PutResultUnit( ResultUnit ru ) throws RemoteException;
}
```


Part of Java-RMI

```
public class WorkUnit implements Serializable {
    double alpha,beta;
    // needed for Serializable interface (needed by RMI)
    private static final long serialVersionUID = 1L;
}
```

Part of Java

```
public class ResultUnit implements Serializable {
    WorkUnit wu;
    double rtnval;
    // needed for Serializable interface (needed by RMI)
    private static final long serialVersionUID = 1L;
}
```

What happens in the Example Code?



Some Caveats

- ✧ The example shown here:
 - ◆ Does NOT provide verification or redundancy in calculations
 - ◆ Does NOT provide redundancy for hardware failures
 - ◆ You could LOSE some results if the network is flaky
 - or if one of the machines is powered off
- ✧ It is suitable for cluster computing environments
 - ◆ Hardware is reasonably stable
 - ◆ Unlikely that there are any “hackers” trying to crack your app
 - ◆ Not good for wide-area or grid computing environments

Building the Client

✱ The Client build is fairly straightforward

- ◆ We need the WorkUnit, ResultUnit, and RmiInterface classes built first (and in that order)

```
jbp@hostname [ 12 ] % make Client.class
javac -classpath ../ WorkUnit.java
javac -classpath ../ ResultUnit.java
javac -classpath ../ RmiInterface.java
javac -classpath ../ Client.java
```

Directory structure:

```
./ClientServer/:
WorkUnit.java
ResultUnit.java
RmiInterface.java
Client.java
Server.java
makefile
```

Building the Server

✱ The Server build is also fairly straightforward ... with one extra step

- ◆ We need the WorkUnit, ResultUnit, and RmiInterface classes built first (and in that order)

```
jbp@hostname [ 22 ] % make Server.class
javac -classpath ../ Server.java
( cd ../ ; rmic ClientServer.Server )
```

↑↑
rmic is the RMI compiler that creates a server "stub" class.
The stub class holds all the networking bits.

↓

```
jbp@hostname [ 36 ] % ls
Client.class      RmiInterface.class  Server_Stub.class
Client.java       RmiInterface.java   WorkUnit.class
ResultUnit.class  Server.class         WorkUnit.java
ResultUnit.java   Server.java          makefile
```

Running the Server

- ✖ Almost always need to start the server first
 - ◆ Unless you've programmed the clients to wait
- ✖ The Naming.rebind() call requires an RMI "Registry"
 - ◆ Luckily, Sun gives you a registry-server with Java

```
jbp@hostname [ 32 ] % rmiregistry &  
jbp@hostname [ 33 ] %
```

put the rmiregistry executable
into the "background"

- ◆ Now start the server

```
jbp@hostname [ 40 ] % java Server.class &  
jbp@hostname [ 41 ] %
```

server URL is hard-coded in executable!

Running the Client(s) ... The Manual Approach

- ✖ Open a bunch of terminals on all of your client machines
 - ◆ In each window, execute the client code:

```
jbp@client1 [ 1 ] % java Client.class 0 6 hostname
```

```
jbp@client2 [ 1 ] % java Client.class 1 6 hostname
```

```
jbp@client3 [ 1 ] % java Client.class 2 6 hostname
```

```
jbp@client4 [ 1 ] % java Client.class 3 6 hostname
```

```
jbp@client5 [ 1 ] % java Client.class 4 6 hostname
```

```
jbp@client6 [ 1 ] % java Client.class 5 6 hostname
```

what client are we?
(informational)

total number of clients
(informational)

server hostname

Running the Client and Server Together

- ✱ ‘rmirun’ is a home-grown Perl script that will launch the registry, one server, and (N-1) clients

```
jbp@hostname [ 1 ] % rmirun -n 8 Server.class Client.class
```

- ◆ It uses ‘ssh’ to connect to the remote clients
 - Assumes that you can ssh to them without a password (ssh-keygen)
- ◆ In a batch/queuing environment, like SGE, you can use ‘qrsh’ instead of ‘ssh’
 - Then rmirun will interact with the batch scheduler and automatically connect to less-loaded machines

```
#$ -cwd -o output.txt  
#$ -pe low* 4-20  
rmirun -S n NSLOTS -m $TMPDIR/machines Server.class Client.class
```

This is home-grown code,
NOT part of Java-RMI

‘rmirun’ and Application Start-Up

- ✱ Recall that we added code to the clients to read some command-line arguments

```
int SelfCID = Integer.parseInt(argv[0]);  
int NumClients = Integer.parseInt(argv[1]);  
String server_url = "/" + argv[2] + "/AlphaBetaServer";
```

- ◆ None of this is really necessary
 - In a stand-alone config (in your own lab), you could hard-code the server URL directly into the client
 - But you’d have to recompile the clients just to move the server
 - In a cluster, the server could be scheduled onto any machine
- ◆ The “MPI Programmer” in me wants to know how many clients we started and which one we are ... not really necessary
 - But it is useful during debugging, since otherwise all client error messages look the same
 - “client-0 error”, “client-1 error”, etc., is much more helpful

High-Performance, Multi-Threaded Server

- ✱ The server can be an obvious bottleneck -- all the clients need to talk to it, and if the server is slowed down then all clients will feel it
 - ◆ Luckily, a Java-RMI server is automatically a multi-threaded server that can take advantage of multiple CPUs on the machine!
- ◆ But the code shown previously is really single-CPU
 - We used the 'synchronized' statement to guarantee correct results
- ◆ One simple approach would be to have several internal `GetNextWorkUnitX()` functions which are independently 'synchronized' ... and key off of client-number to send some clients to each X
 - Since they are not globally synchronized, 1 CPU can be inside of each `X()` function ... meaning you have a multi-CPU server!

Multi-Threaded Server Example

```
public interface RmiInterface extends Remote {
    public WorkUnit GetNextWorkUnit( int clnum ) throws RemoteException;
    public int PutResultUnit( ResultUnit ru ) throws RemoteException;
}
```

```
...
class Server extends UnicastRemoteObject implements RmiInterface {
    ...
    public synchronized WorkUnit GetNextWorkUnitEven() throws RemoteException {
        WorkUnit wu = new WorkUnit();
        wu.alpha = next_alpha;
        wu.beta = next_beta;
        next_alpha += 2.0*delta;
        next_beta -= 2.0*delta;
        return( wu );
    }
    public synchronized WorkUnit GetNextWorkUnitOdd() throws RemoteException {
        WorkUnit wu = new WorkUnit();
        wu.alpha = next_alpha + delta;
        wu.beta = next_beta + delta;
        next_alpha += 2.0*delta;
        next_beta -= 2.0*delta;
        return( wu );
    }
    public WorkUnit GetNextWorkUnit( int clnum ) throws RemoteException {
        if( (clnum%2) == 0 ) {
            wu = GetNextWorkUnitEven();
        } else {
            wu = GetNextWorkUnitOdd();
        }
        return( wu );
    }
}
```

Handling Errors

- ✖ While modern networks are fairly robust, it is always a good idea to detect network errors (RemoteException errors) and do something “intelligent” with them
 - ◆ Often, you can just retry the network and it will work fine
- ✖ If you are doing wide-area (even cross-campus) networking, then you are **LIKELY** to see some errors ... on both the client and server
 - ◆ Again, many of them will go away with a retry
- ✖ Note that some network connection routines have long “time-outs”
 - ◆ Sometimes a function will wait as much as 5 minutes before saying the network has problems!

Network Errors, Client

- ✖ UnknownHostException (Naming.lookup)
 - ◆ Yes, it’s the obvious cause
- ✖ ConnectException (Naming.lookup or server.function)
 - ◆ The connection was refused by the remote machine, i.e. the network is ok, but the server is refusing to hand out any data
 - ◆ Or a network time-out or other network problem ... retry?
- ✖ ConnectIOException
 - ◆ Possible network problems ... retry?
- ✖ MarshalException (server.function)
- ✖ UnmarshalException (server.function)
 - ◆ Problems when “marshalling” the data to the remote machine ... possibly just a problem with the Serialization of your data-type; but could also be an RMI-version mismatch
 - ◆ This is unlikely, but very bad when it happens (the data **MIGHT** have gotten to the server and been processed, but you can’t tell)

Network Errors, Client, cont'd

- ✖ `ServerError` (`server.function`)
- ✖ `ServerException` (`server.function`)
- ✖ `ServerRuntimeException` (`server.function`)
 - ◆ The server encountered an error in processing the remote method ... could be bad parameters were sent; could be that the server code is bad ... retry?
 - which might lead to ...
- ✖ `NoSuchObjectException` (`server.function`)
 - ◆ The server is no longer running ... restart the server and try again
 - But “restart the server” is an off-line process

Network Errors, Server

- ✖ Aside from the network-badness errors (just like the client might see), the server might also see:
- ✖ `AlreadyBoundException` (`Naming.rebind`)
 - ◆ Someone else is using the same server-name ... so try another?
 - But your client is probably expecting the server at a given URL
 - ◆ If you're sharing the server machine with others (like in the DSCR), and you've shared your code with others, then someone else beat you to that service name
 - Might be worth thinking about NOT hard-coding the name into your clients or server (yet another command-line argument?)

Results Checking

- ✖ If you are doing wide-area (even cross-campus) networking, then you are **LIKELY** to see some weird data hit your server
 - ◆ Some of this could be hackers who are actively trying to screw your program up
 - ◆ Lots of it will just be security scanners that are sending random packets out to every machine to see what happens
 - Either way, your server should really do some verification
 - ◆ Some of this garbage-data may look like real results
 - So you may want to run every WorkUnit twice and make sure both results match
 - ◆ Given all the weird errors that can creep in, you may want to allow clients to re-send a result several times, to make sure it gets through

Results Checking, cont'd

- ✖ If you are doing wide-area (even cross-campus) networking, then you are **ALSO** likely to see some client slow-downs
 - ◆ Might want to track the date/time that a job was handed to a client
 - If it takes too long to get back a result, re-compute that WorkUnit by sending it to another client
- ✖ So we're now looking at a system which submits each WorkUnit to several different clients, with each client returning multiple ResultUnit packets for redundancy, while checking for garbage data, black-listing hackers, but not erroneously labeling slow clients as hackers
 - Piece of Cake!! No Problem!!
- ◆ Store current and completed work-units in a HashSet, List, or array

Famous Last Words

- ✧ If you're running the client and server (both) on a cluster, then you probably don't have to worry much about network errors or bad results packets from hackers
 - ◆ Probably a good idea to still catch the errors and print them, just for debugging
 - Actually, Java requires that you catch them

Famous Last Words, cont'd

- ✧ On a cluster, you are still sharing machines with others ...
- ✧ What if other people are using your cool, new, super-fast, parallel client-server app?
 - ◆ They're all going to start the same named servers ... bad!
 - You may want to have a scheme for re-naming the server if the "usual" name is already in use
 - Prepend your NetID or \$JOBID; append a '1', then '2', ...
- ✧ On the DSCR, a low-priority client can be dropped to 1% CPU when a high-priority job gets started on the same machine
 - ◆ A 100x slow-down is significant!
 - ◆ You may want the server to detect slow-downs and issue redundant WorkUnits
 - But you can probably assume that the first result back is correct

Next Steps

- ✱ Add a database back-end for job monitoring
 - ◆ Using a HashSet, List, or array means storing data in memory
 - So you could run out of memory on the server machine
 - ◆ Store the data to a database (file)
 - Java has support for using those database files as if they were “regular” HashSet, Maps, Lists, etc.
 - ◆ JDBC+SQLite
 - JDBC is a set of connector-functions to a “real” database
 - SQLite is a mostly-full-featured SQL server
 - Used in many production apps (Mac OS X)
 - Database will store data to disk
 - So server failure won’t be catastrophic

“Message Passing” with Java-RMI ??

- ✱ Java-RMI is more broad than just client-server
 - ◆ One could imagine each “client” also acting as a “server”

```
public interface RmiInterface extends Remote {
    public MessageUnit GetMessage() throws RemoteException;
    public int PutMessage( int other_cli, MessageUnit mu ) throws RemoteException;
}
```

